



# LAB 3: FULLY CONNECTED DEEP NETWORKS (FCN)

University of Washington, Seattle

Spring 2022



# OUTLINE

## [Part 1: Introduction to FCN](#) ([video link](#))

- [Shallow vs Deep Networks](#)
- [Regression vs Classification](#)
- [Outputting Probabilities with Softmax Function](#)

## [Part 2: Additional Data Prep Methods](#) ([video link](#))

- [Training/Validation/Test sets](#)
- [One-hot encoding](#)

## [Part 3: Stochastic Gradient Descent](#) ([video link](#))

- [SGD, Mini-batch GD, Batch GD](#)
- [Implementing variants of Gradient Descent](#)

## [Part 4: Additional Hyperparameters & Regularizers](#) ([video link](#))

- [Activation functions](#)
- [Loss functions and Advanced Optimizers](#)
- [Regularizers](#)
- [Batch Normalization](#)
- [Weight initializations](#)

## [Part 5: Training Fully Connected Deep Network](#) ([video link](#))

- [Iris Classification Example](#)

## [Lab Assignment](#) ([video link](#))

- [MNIST Classification using FCN](#)
- [Tips for Training Your Model](#)



# INTRODUCTION TO FCN

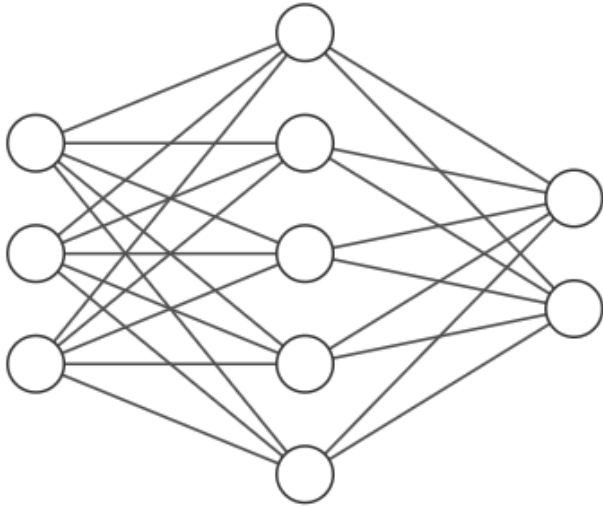
Shallow vs Deep Networks

Regression vs Classification

Outputting Probabilities with Softmax Function



# Shallow vs Deep Networks



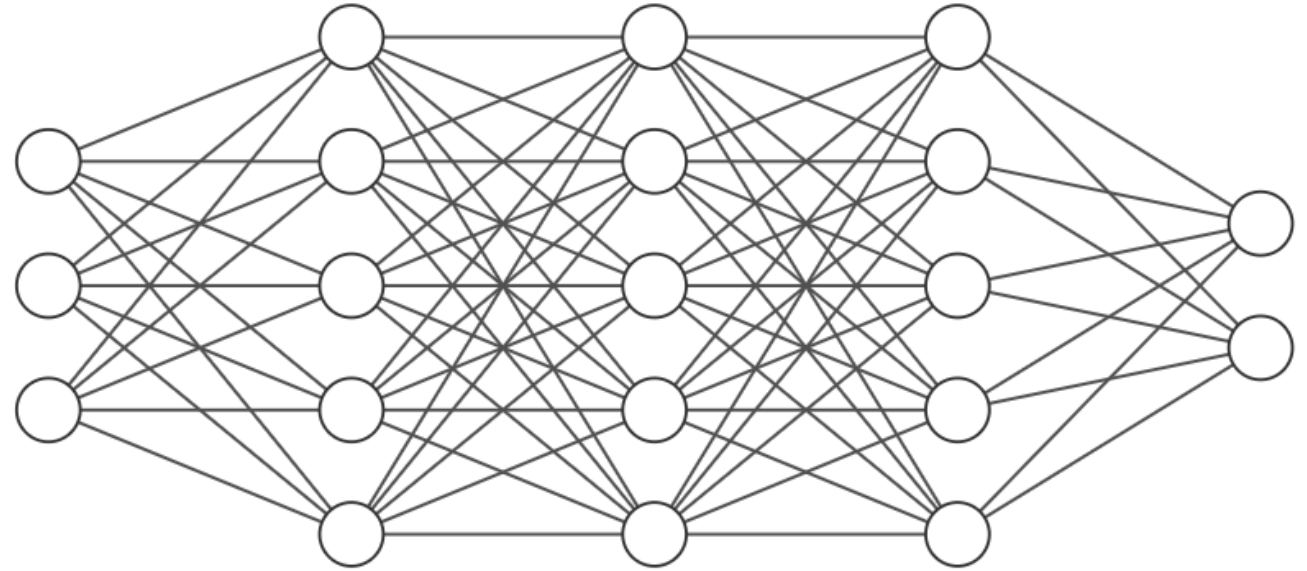
## Shallow Networks

One hidden layer

Universal function approximator  
(with enough hidden neurons)

Easier to train

Fit for simple problems



## Deep Networks

>1 hidden layer

Upper hidden neurons reuse lower-level features  
Can approximate more complex & general functions

Harder to train

Fit for more complex problems



# Regression vs Classification

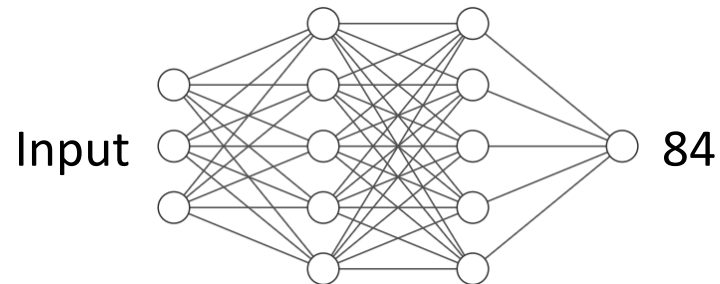
## Regression



What will be the temperature tomorrow?



Fahrenheit



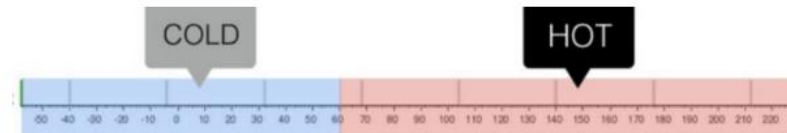
Input

84

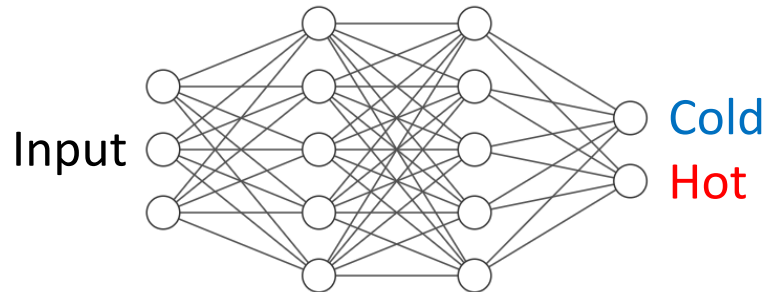
## Classification



Will it be hot or cold tomorrow?



Fahrenheit

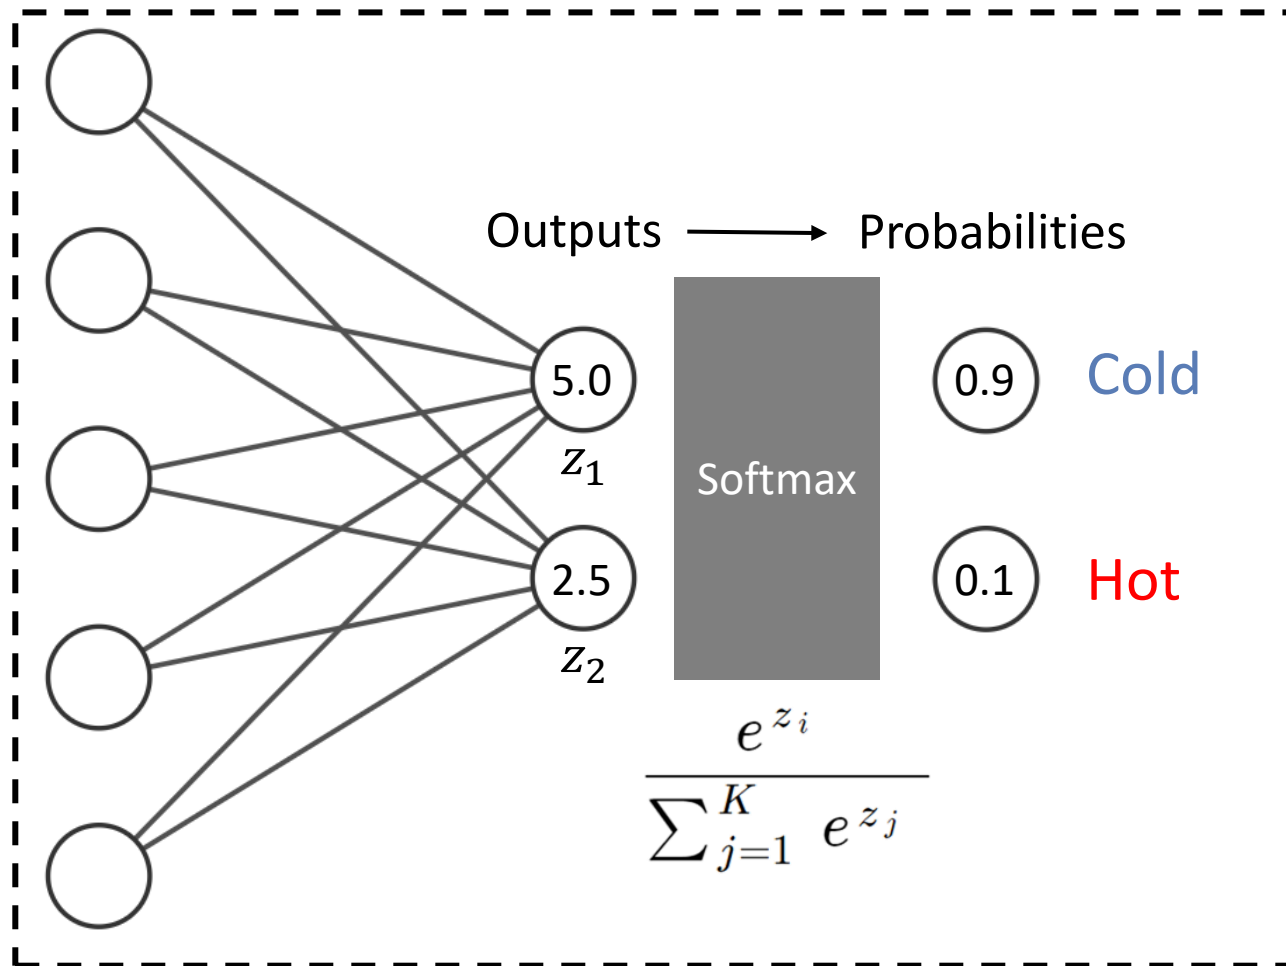
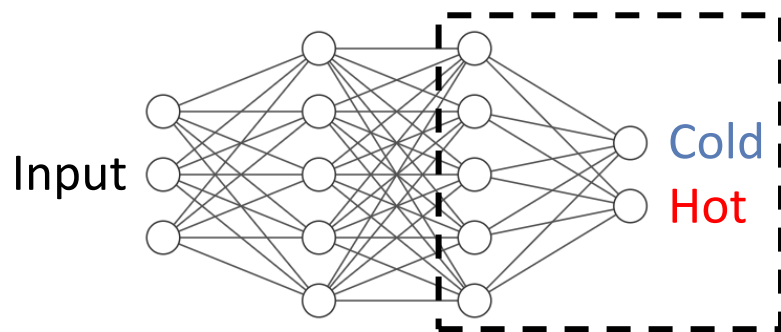


Input

Cold  
Hot

Image credit: Towards Data Science

# Outputting Probabilities with Softmax Function



`torch.nn.Softmax()`



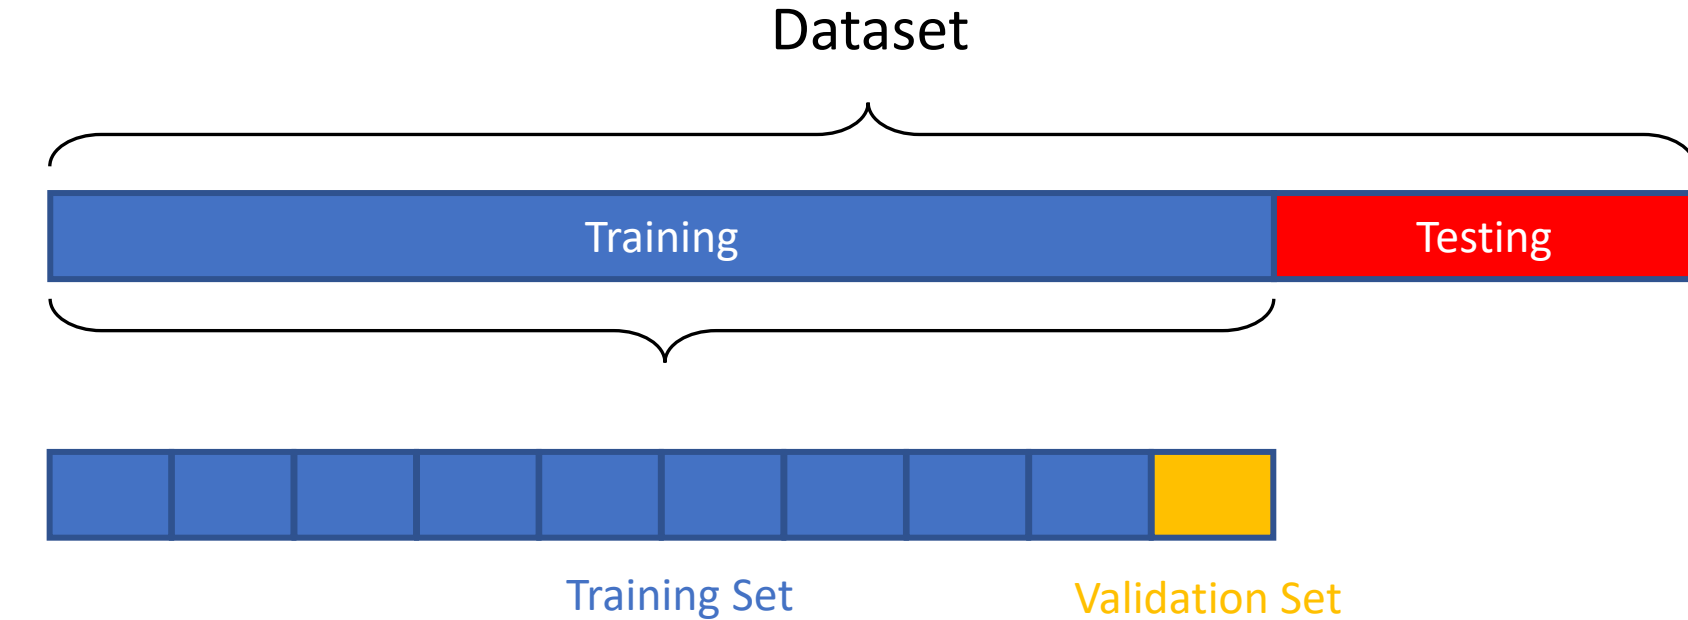
# ADDITIONAL DATA PREP METHODS

Training/Validation/Test Sets

One-Hot Encoding



# Train/Validation/Test Split



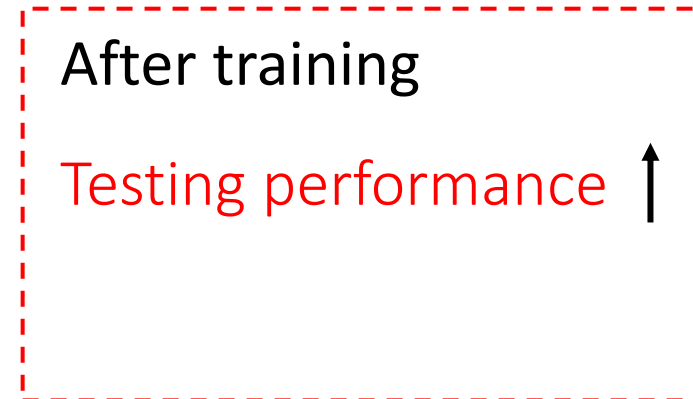
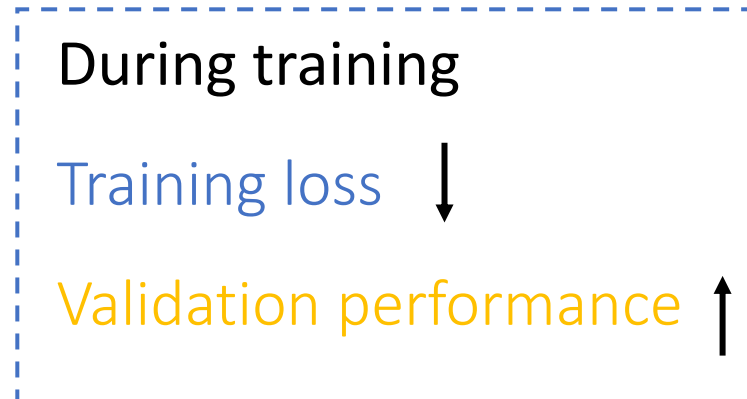
Common ratios

70, 15, 15

80, 10, 10

60, 20, 20

⋮







# Train/Validation/Test Split

```
1 from sklearn.datasets import load_iris
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.model_selection import train_test_split
4
5 iris = load_iris()
6
7 X = iris['data']
8 y = iris['target']
9
10 scaler = StandardScaler()
11 X_scaled = scaler.fit_transform(X)
12
13 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
14                                                    Testing data ratio (0.2 = 20%) ← test_size=0.2,
15                                                    Random seed to use for splitting ← random_state=2)
16
17 X_validation = X_train[:int(len(X_test))]
18 y_validation = y_train[:int(len(X_test))]
19
20 X_train = X_train[int(len(X_test)):]
21 y_train = y_train[int(len(X_test)):]
```

Load Iris dataset

Extract features (X) and target labels (y)

Scale features using scikit-learn provided standard scaler

Split the dataset into Training (80%) and Testing (20%)

Assign subset of the training dataset as validation data (same size as testing)

Use the remaining dataset as training

Final split ratios = Training: 60%, Testing: 20%, Validation: 20%



# One-hot Encoding

Human Readable 

Dog →

Cat →

Turtle →

Fish →

Snake →

Categorical Data

Machine Readable 

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 1 | 0 | 0 | 0 |
|---|---|---|---|---|

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 0 | 1 | 0 | 0 |
|---|---|---|---|---|

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 0 | 0 | 1 | 0 |
|---|---|---|---|---|

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 1 |
|---|---|---|---|---|

One-hot vectors



# One-hot Encoding

```
1 labels = torch.tensor([0,1,2])  
2 print(labels)
```

Original labels are 0, 1 and 2

```
tensor([0, 1, 2])
```

```
1 labels_one_hot = torch.nn.functional.one_hot(labels)  
2 print(labels_one_hot)
```

```
tensor([[1, 0, 0],  
        [0, 1, 0],  
        [0, 0, 1]])
```

Labels after one-hot encoding  
(torch.nn.functional.one\_hot())

Official function documentation:

[https://pytorch.org/docs/stable/generated/torch.nn.functional.one\\_hot.html](https://pytorch.org/docs/stable/generated/torch.nn.functional.one_hot.html)



# STOCHASTIC GRADIENT DESCENT

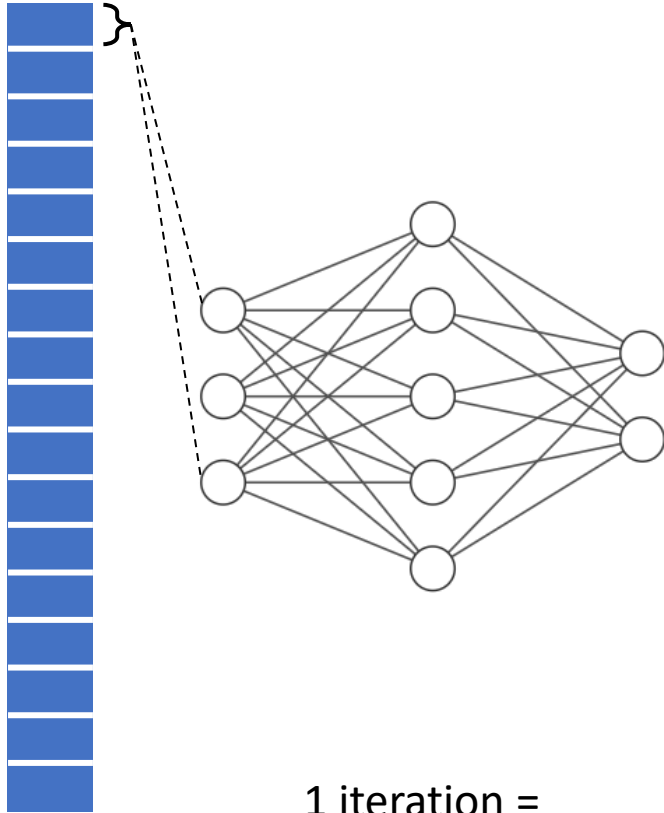
SGD, Mini-batch GD, Batch GD



# Variants of Gradient Descent

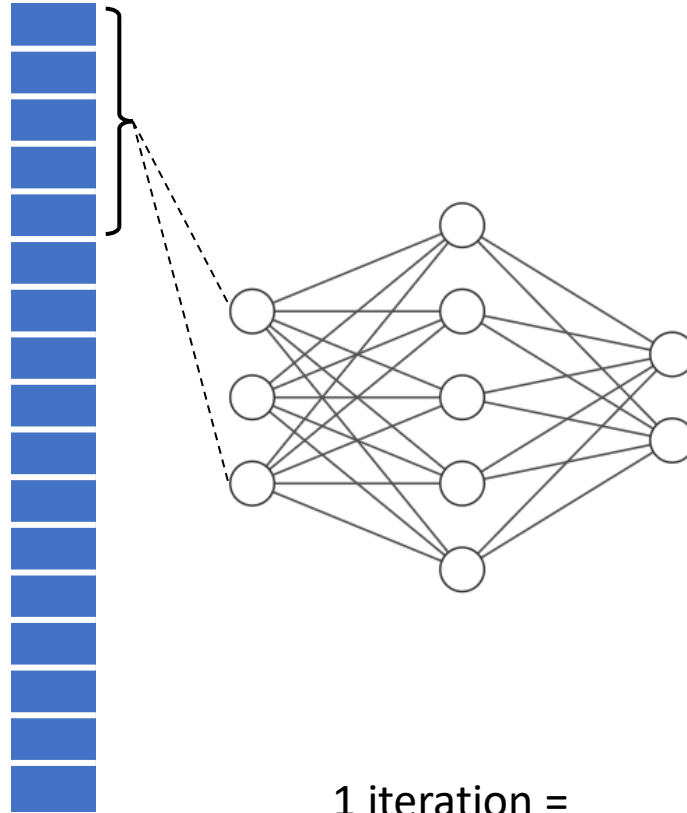
Training  
Dataset

SGD



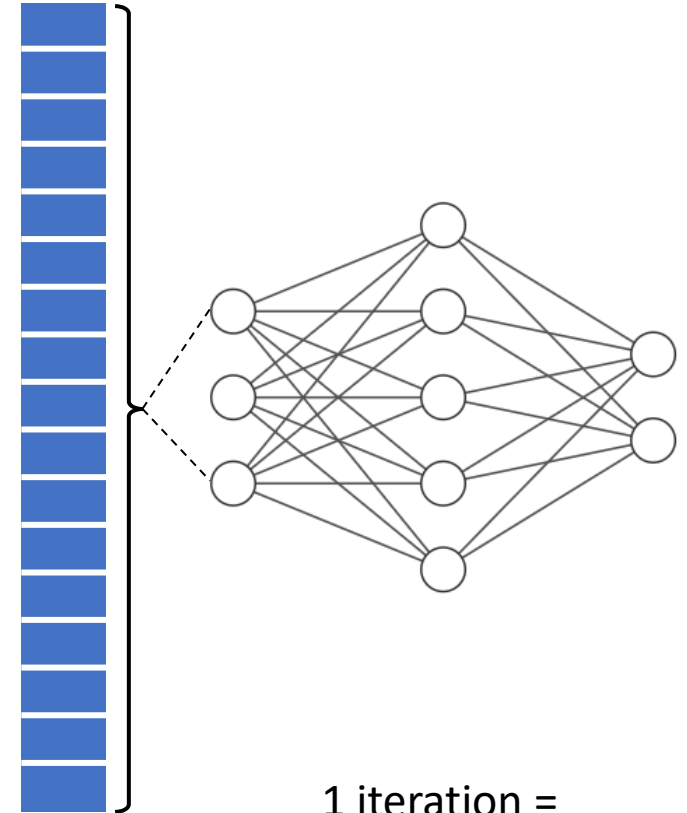
1 iteration =  
Fwd/bwd pass 1 training sample

Mini-batch



1 iteration =  
Fwd/bwd pass n-training samples  
( $n < \text{total \# of training samples}$ )

Batch GD



1 iteration =  
Fwd/bwd pass full training samples

# Implementing Variants of Gradient Descent

## SGD

For **epoch** in range(**epochs**)

For sample in train

### Training loop

- zero\_grad
- fwd pass input
- compute loss
- backpropagation
- update weights/biases

Total # of iterations      Epochs  $\times$  m  
(m = total # of samples in training)

## Mini-batch

For **epoch** in range(**epochs**)

For mini-batch in train

### Training Loop

Epochs  $\times$  R  
(R = m/mini-batch size)

## Batch GD

For **epoch** in range(**epochs**)

### Training Loop

Epochs



# ADDITIONAL HYPER-PARAMETERS & REGULARIZERS

Activation Functions

Loss Functions and Advanced Optimizers

Regularizers

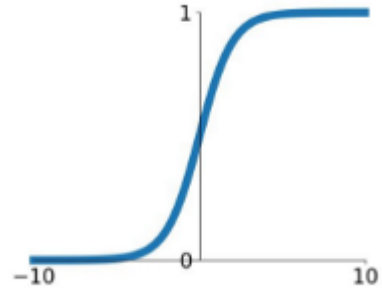
Batch Normalization

Network Initialization



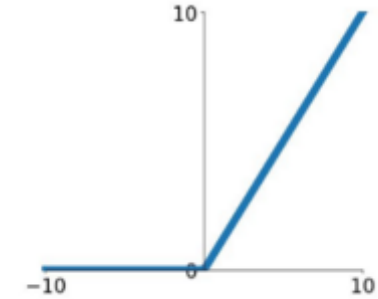
# Activation Functions

Sigmoid



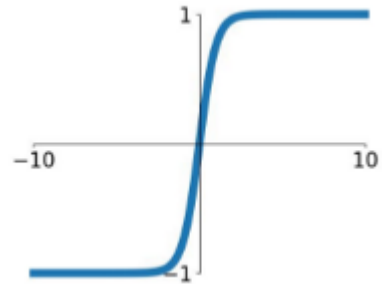
`torch.nn.functional.sigmoid()`

ReLU



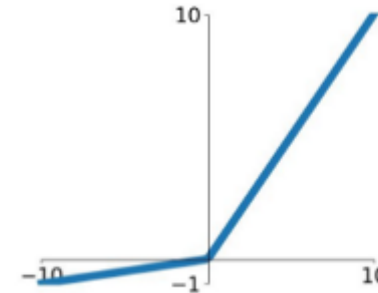
`torch.nn.functional.relu()`

Tanh



`torch.nn.functional.tanh()`

Leaky ReLU



`torch.nn.functional.leaky_relu()`





# Activation Functions

```
1 class Model(torch.nn.Module):
2
3     def __init__(self, input_dim, output_dim):
4
5         super(Model, self).__init__()
6
7         self.layer1 = torch.nn.Linear(input_dim, 5)
8         self.layer2 = torch.nn.Linear(5, 5)
9         self.layer3 = torch.nn.Linear(5, output_dim)
10
11     def forward(self, x):
12
13         [out1 = torch.nn.functional.relu(self.layer1(x))
14         out2 = torch.nn.functional.relu(self.layer2(out1))
15         output = torch.nn.functional.softmax(self.layer3(out2), dim=1)
16
17     return output
```

Apply ReLU activation function to outputs of layer1 and 2



# Loss Functions

## Regression

Mean Squared Error  $\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$  `torch.nn.MSELoss()`

## Classification

Cross Entropy  $L_{\text{CE}} = - \sum_{i=1}^n t_i \log(p_i)$  `torch.nn.CrossEntropyLoss()`

NOTE: `torch.nn.CrossEntropyLoss()` automatically implements **one-hot encoding** and **softmax** when providing integer labels as targets.



# Advanced Optimizers

Gradient Descent

`torch.optim.SGD()`

Adam

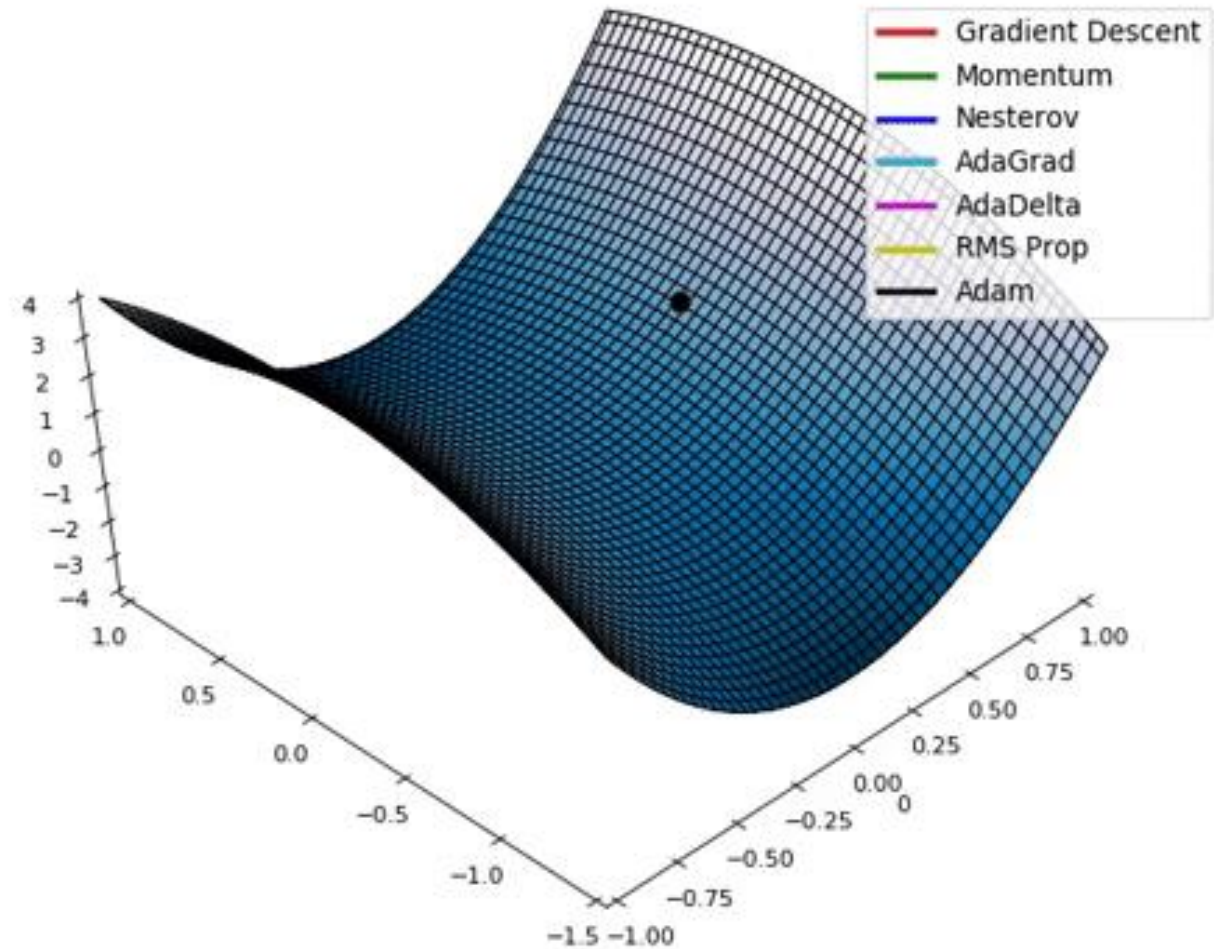
`torch.optim.Adam()`

RMS Prop

`torch.optim.RMSprop()`

Ada Delta

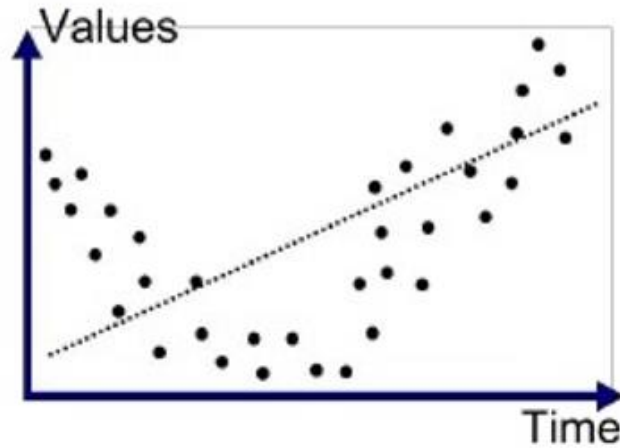
`torch.optim.Adadelta()`



More Optimizers: <https://pytorch.org/docs/stable/optim.html>

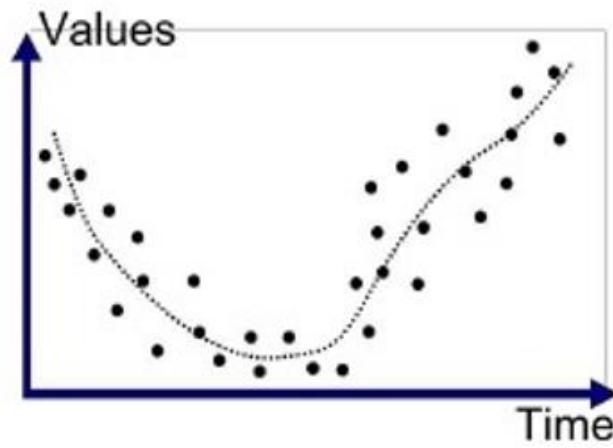


# Avoiding Overfitting with Regularization



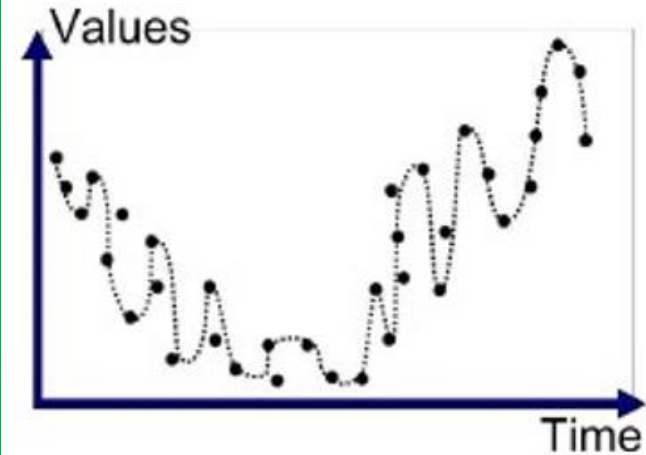
Underfitted

Bad training accuracy  
Bad testing accuracy



Good Fit/Robust

Good training accuracy  
Good testing accuracy



Overfitted

Great training accuracy  
Bad testing accuracy



# L1, L2 Regularizations in PyTorch

## L1 Regularization

$$Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N |w_i|$$

Penalizes **sum of absolute values of weights**

Results in a sparse model

Not suitable for learning complex patterns

Robust to outliers

## L2 Regularization

$$Loss = Error(y, \hat{y}) + \lambda \sum_{i=1}^N w_i^2$$

Penalizes **sum of squared values of weights**

Results in a dense model

Learns complex patterns

Sensitive to outliers



# L1, L2 Regularizations in PyTorch

## L1

```
1 l1_penalty = torch.nn.L1Loss(size_average=False)
2 reg_loss = 0
3
4 for param in model.parameters():
5
6     reg_loss += l1_penalty(param)
7
8 lambda_ = 0.9
9
10 loss += lambda_ * reg_loss
```

Loop through model parameters to compute L1 regularization term

Pick the lambda value

Add the L1 term to loss during training

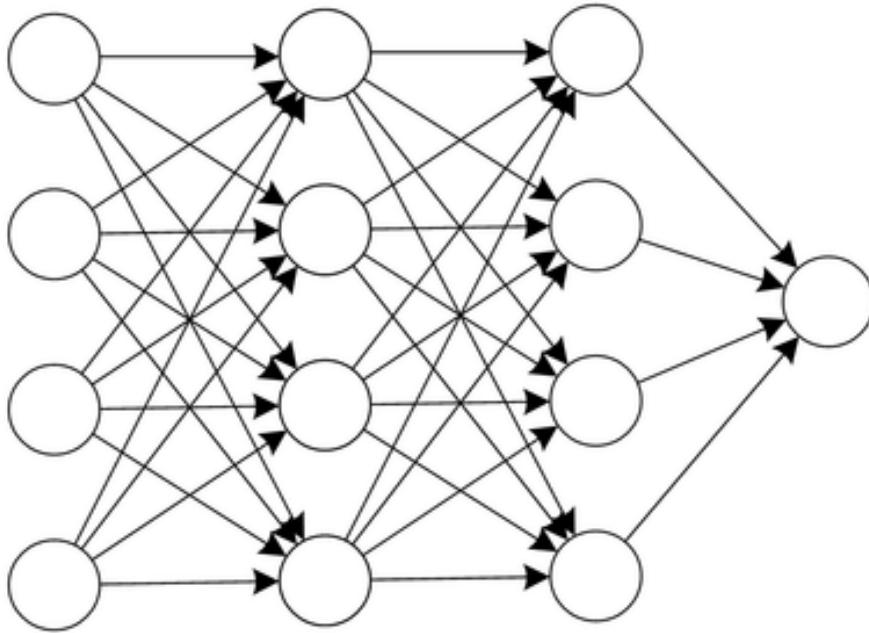
## L2

```
1 optimizer = torch.optim.SGD(model.parameters(), lr=0.001, weight_decay=0.9)
```

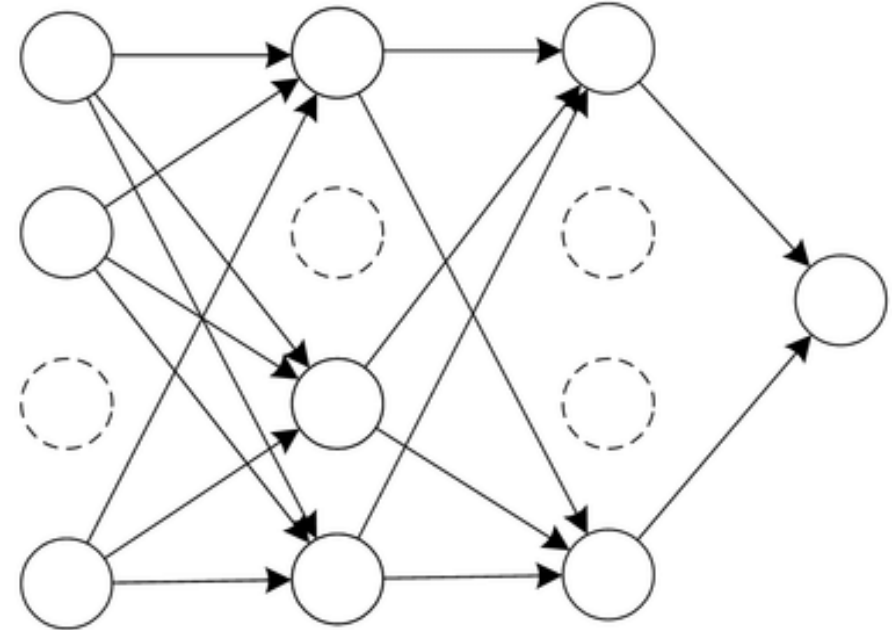
weight\_decay sets lambda value for L2 regularization term



# Dropout Regularization in PyTorch



Standard Neural Network



Network with Dropout

[Srivastava et al 2014 \(~35000 citations!\)](#)

Dropout forces the network to learn **more robust features** +  
**different random subsets** of other neurons



# Dropout Regularization in PyTorch

```
1 class Model(torch.nn.Module):
2
3     def __init__(self, input_dim, output_dim):
4
5         super(Model, self).__init__()
6
7         self.layer1 = torch.nn.Linear(input_dim, 5)
8         self.layer2 = torch.nn.Linear(5, output_dim)
9         self.dropout = torch.nn.Dropout(p = 0.25)
10
11     def forward(self, x):
12
13         out1 = self.layer1(x)
14         out1 = self.dropout(out1)
15         output = self.layer2(out1)
16
17         return output
```

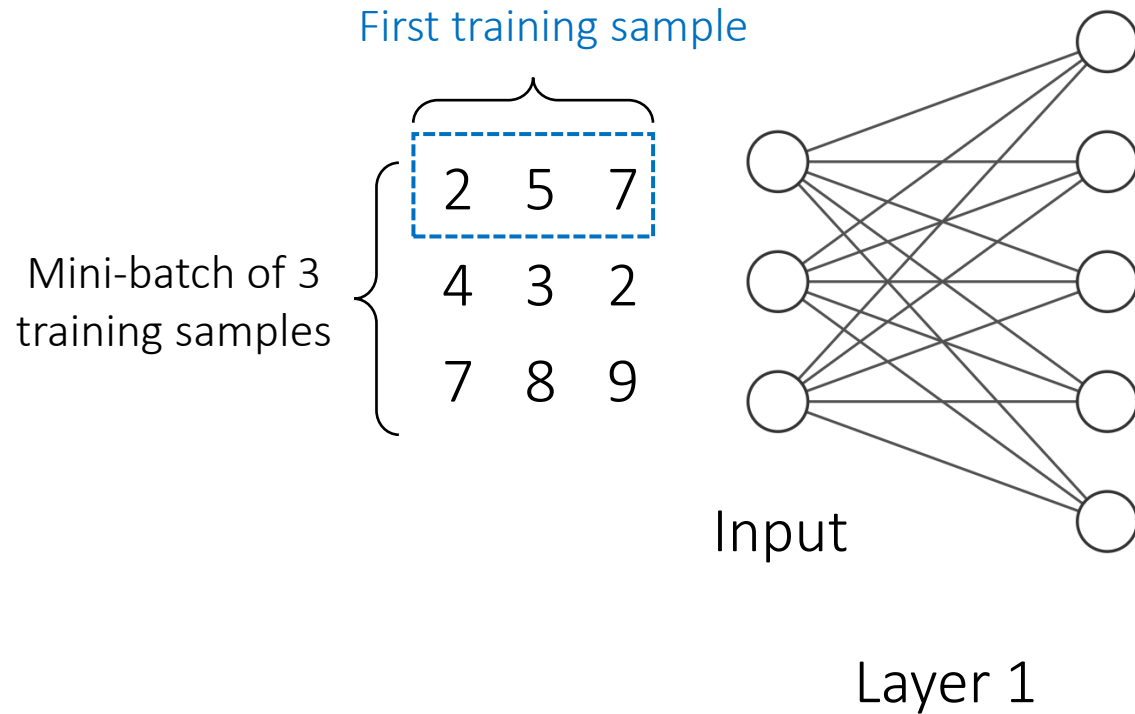
Add the dropout in `__init__` with  
`p` = probability of neuron state is set to 0

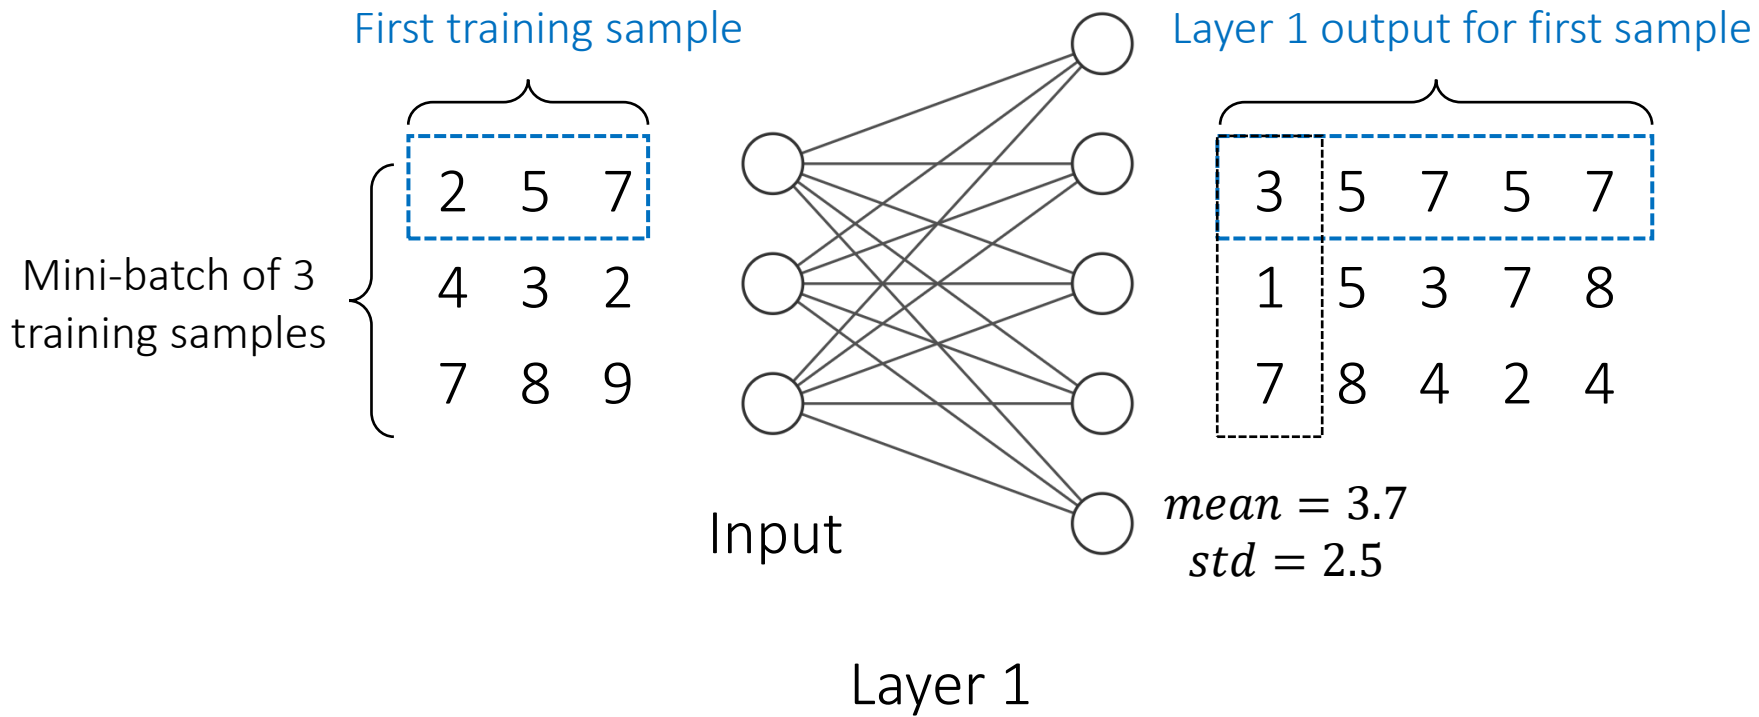
Apply dropout to the output of the  
desired layer

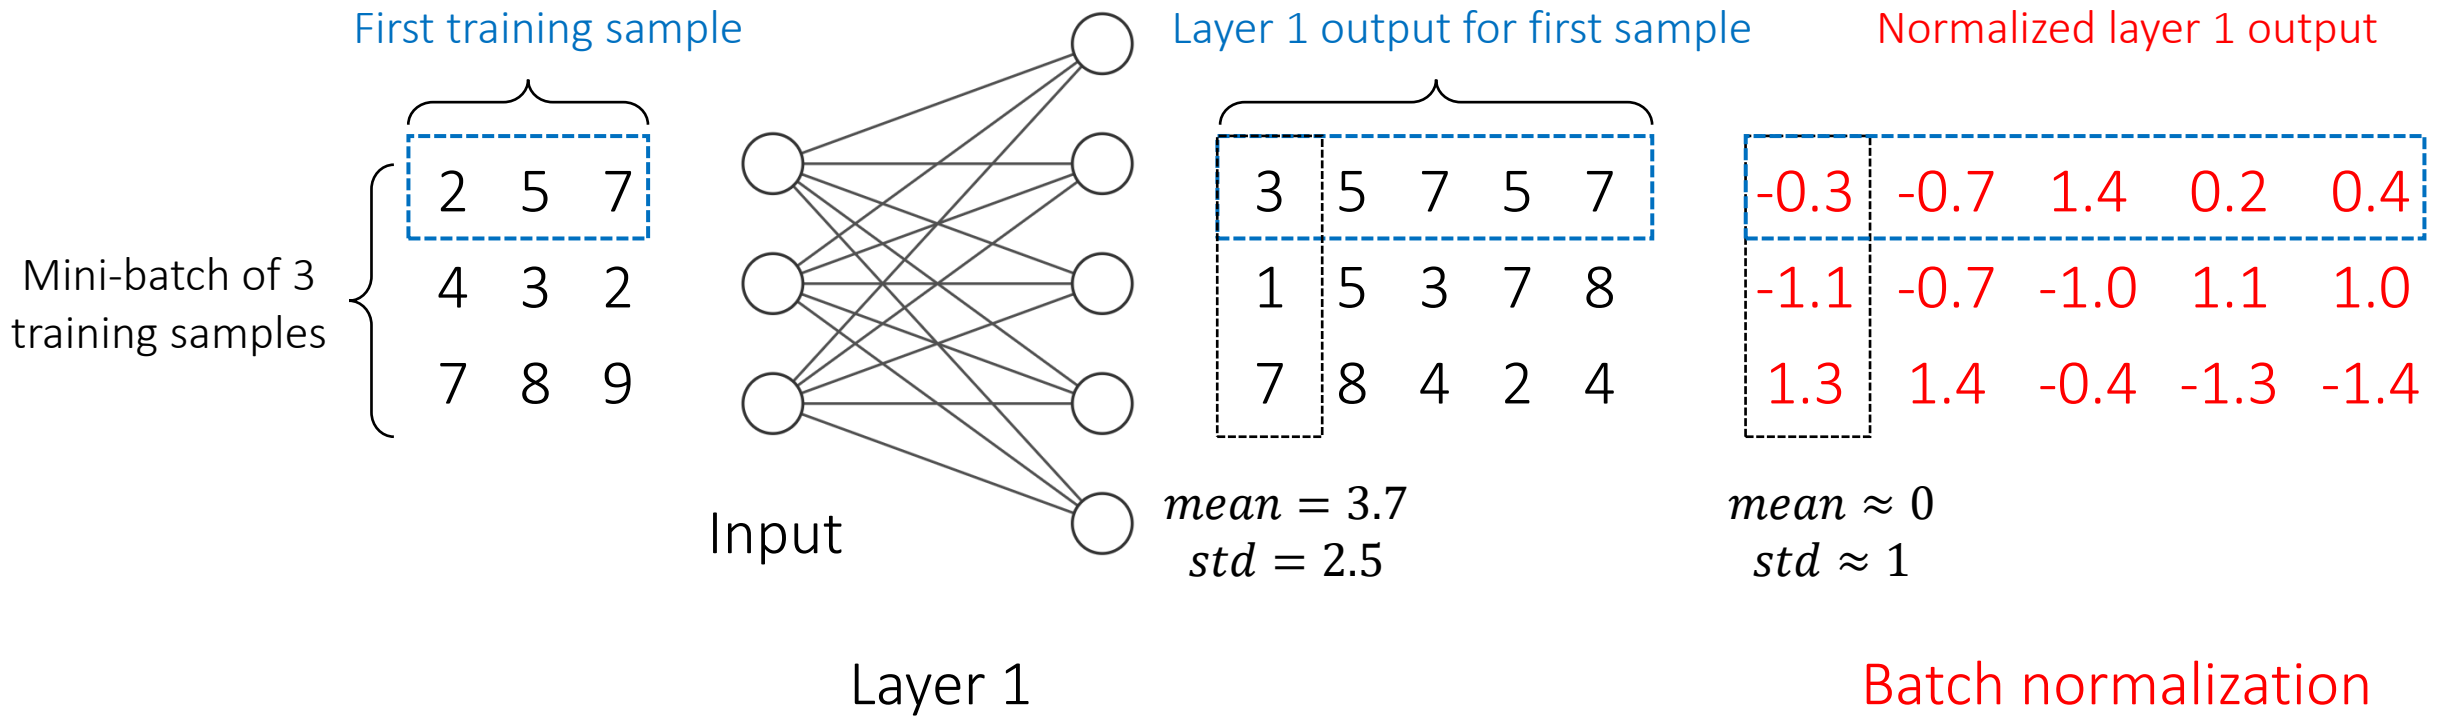


# Preventing Vanishing/Exploding Gradients: Batch Normalization

[Lofte et al 2015 \(>35000 citations\)](#)









# Batch Normalization in PyTorch

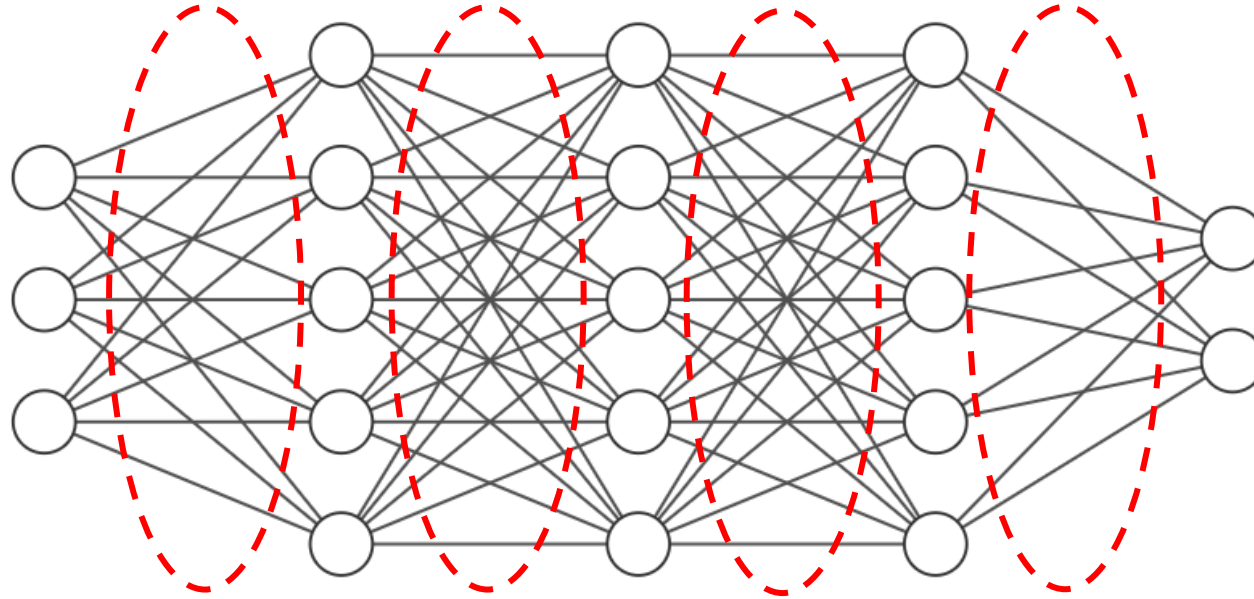
```
1 class Model(torch.nn.Module):
2
3     def __init__(self, input_dim, output_dim):
4
5         super(Model, self).__init__()
6
7         self.layer1 = torch.nn.Linear(input_dim, 5)
8         self.layer2 = torch.nn.Linear(5, output_dim)
9         self.bn1 = torch.nn.BatchNorm1d(5)
10
11     def forward(self, x):
12
13         out1 = self.layer1(x)
14         out1 = self.bn1(out1)
15         out1 = torch.nn.functional.relu(out1)
16         output = self.layer2(out1)
17
18     return output
```

Add BatchNorm1D in `__init__` with number of features in desired layer output

Apply batch normalization to the output of the desired layer.

**NOTE:** Batch normalization is done **BEFORE** feeding into activation function

# Preventing Vanishing/Exploding Gradients: Weight Initialization



Proper weight initialization plays essential roles in preventing exploding/vanishing gradients



Faster convergence



# Weight Initialization Methods

## Uniform

Uniform distribution

`torch.nn.init.uniform_()`

## Normal

Gaussian distribution

`torch.nn.init.normal_()`

## Xavier

Suitable for `tanh()` activation

`torch.nn.init.xavier_uniform_()`

[Xavier et al 2010](#)

## Kaiming

Suitable for `ReLU()` activation

`torch.nn.init.kaiming_uniform_()`

[He et al 2015](#)

More initializations: <https://pytorch.org/docs/stable/nn.init.html>



# Weight Initialization in PyTorch

```
1 class Model(torch.nn.Module):
2
3     def __init__(self, input_dim, output_dim):
4
5         super(Model, self).__init__()
6
7         self.layer1 = torch.nn.Linear(input_dim, 5)
8         self.layer2 = torch.nn.Linear(5, output_dim)
9         [torch.nn.init.kaiming_uniform_(self.layer2.weight)]
10
11     def forward(self, x):
12
13         out1 = torch.nn.functional.relu(self.layer1(x))
14         output = torch.nn.functional.relu(self.layer2(out1))
15
16         return output
```

Manually apply Kaiming He Initialization to layer2

**NOTE:** PyTorch already applies good default initialization for most layers. We suggest using manual initialization for experimental purposes



# TRAINING FULLY CONNECTED NETWORKS

Iris Classification Example





# Neural Network Workflow in PyTorch

Prepare Data

Define Model

Select Hyperparameter

Identify Tracked Values

Train Model

Visualization and Evaluation



# Prepare Data

```
1 from sklearn.datasets import load_iris
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.model_selection import train_test_split
4
5 iris = load_iris()
6
7 X = iris['data']
8 y = iris['target']
9
10 scaler = StandardScaler()
11 X_scaled = scaler.fit_transform(X)
12
13 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
14                                                    Testing data ratio (0.2 = 20%) ← test_size=0.2,
15                                                    Random seed to use for splitting ← random_state=2)
16
17 X_validation = X_train[:int(len(X_test))]
18 y_validation = y_train[:int(len(X_test))]
19
20 X_train = X_train[int(len(X_test)):]
21 y_train = y_train[int(len(X_test)):]
```

Load Iris dataset

Extract features (X) and target labels (y)

Scale features using scikit-learn provided standard scaler

Split the dataset into Training (80%) and Testing (20%)

Assign subset of the training dataset as validation data (same size as testing)

Use the remaining dataset as training

Final split ratios = Training: 60%, Testing: 20%, Validation: 20%



# Define Model

```
1 class irisClassificationFCN(torch.nn.Module):
2
3     def __init__(self, input_dim, output_dim, hidden1_dim, hidden2_dim):
4
5         super(irisClassificationFCN, self).__init__()
6
7         self.layer1 = torch.nn.Linear(input_dim, hidden1_dim)
8         self.layer2 = torch.nn.Linear(hidden1_dim, hidden2_dim)
9         self.layer3 = torch.nn.Linear(hidden2_dim, output_dim)
10
11     def forward(self, x):
12
13         out1 = torch.nn.functional.relu(self.layer1(x))
14         out2 = torch.nn.functional.relu(self.layer2(out1))
15         output = self.layer3(out2)
16
17     return output
```

Input layer (input\_dim = 4)

2 hidden layers (hidden1\_dim = 30, hidden2\_dim = 10)

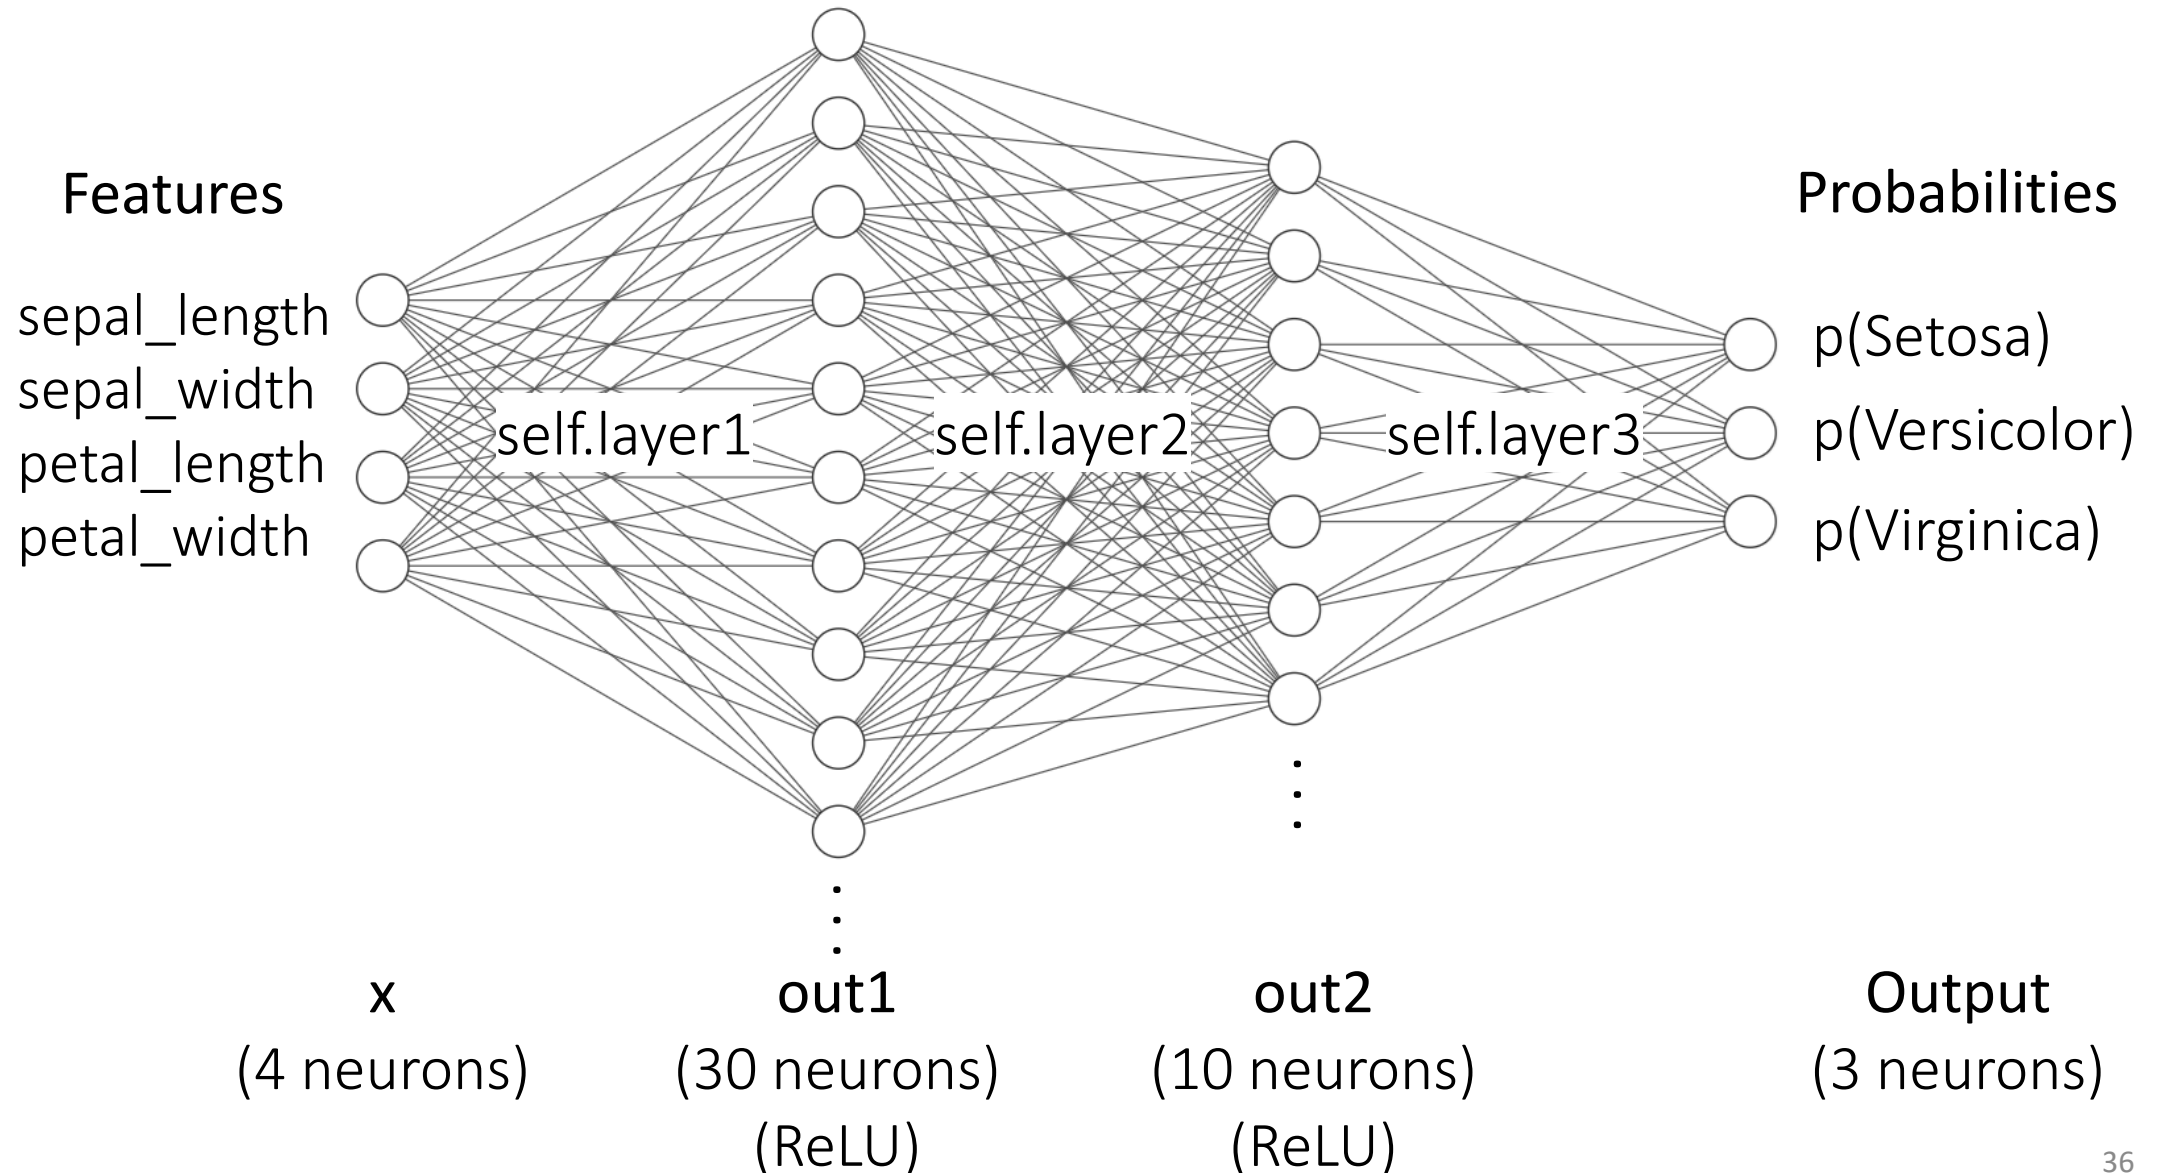
output layer (output\_dim = 3)

ReLU activation for the outputs of each hidden layer

Return raw final output



# Define Model





# Select Hyperparameter

```
1 model = irisClassificationFCN(input_dim = 4, output_dim = 3,  
2                               hidden1_dim = 30, hidden2_dim = 10)  
3  
4 learning_rate = 0.025  
5 epochs = 25  
6  
7 loss_func = torch.nn.CrossEntropyLoss()  
8 optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)  
9  
10 model
```

```
irisClassificationFCN(  
  (layer1): Linear(in_features=4, out_features=30, bias=True)  
  (layer2): Linear(in_features=30, out_features=10, bias=True)  
  (layer3): Linear(in_features=10, out_features=3, bias=True)  
)
```

Initialize the model with input\_dim = 4, output\_dim = 3, hidden1\_dim = 30, hidden2\_dim = 10

Using learning rate of 0.025 and 25 as epochs

Using Cross Entropy Loss and Adam Optimizer

Model variable correctly shows our network structure



# Identify Tracked Values

```
1 train_loss_list = np.zeros((epochs,))  
2 validation_accuracy_list = np.zeros((epochs,))
```

Create empty list or NumPy arrays to hold training loss and validation accuracy



# Train Model

```
1 import tqdm
2
3 train_inputs = torch.from_numpy(X_train).float()
4 train_targets = torch.from_numpy(y_train).long()
5
6 validation_inputs = torch.from_numpy(X_validation).float()
7 validation_targets = torch.from_numpy(y_validation).long()
8
9 testing_inputs = torch.from_numpy(X_test).float()
10 testing_targets = torch.from_numpy(y_test).long()
11
```

Import tqdm to visualize the progress of your training

Convert training/validation/testing datasets into PyTorch Tensors

Convert the targets into int64 form via .long()



# Train Model

```
14 # Training Loop -----
15
16 for epoch in tqdm.trange(epochs):
17     optimizer.zero_grad()
18     train_outputs = model(train_inputs)
19     loss = loss_func(train_outputs, train_targets)
20     train_loss_list[epoch] = loss.item()
21     loss.backward()
22     optimizer.step()
23
24 # Compute Validation Accuracy -----
25
26 with torch.no_grad():
27     validation_outputs = model(validation_inputs)
28     correct = (torch.argmax(validation_outputs, dim=1) ==
29               validation_targets).type(torch.FloatTensor)
30     validation_accuracy_list[epoch] = correct.mean()
```

## Training loop

- Empty gradient buffer
- Forward propagation
- Compute loss
- Save loss value to a list
- Backward propagation
- Update weights/biases

## Compute Validation accuracy per epoch

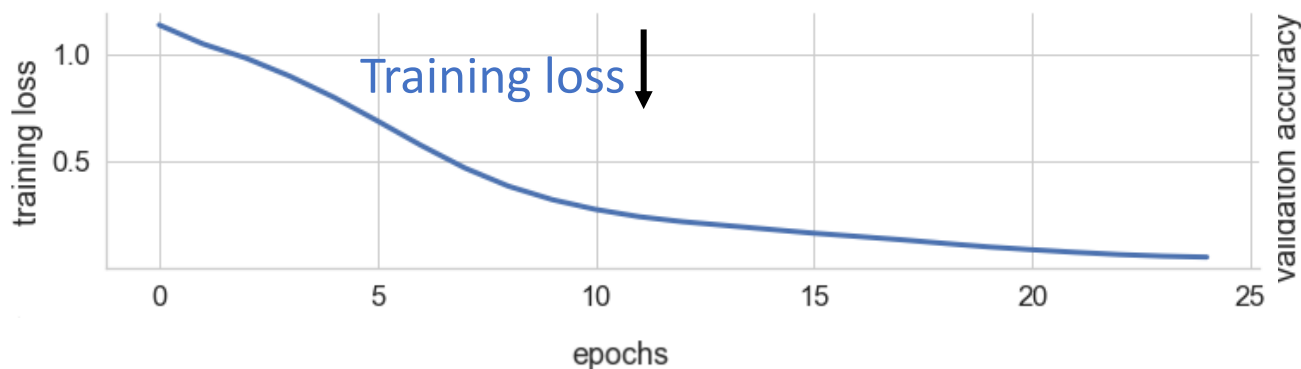
- Forward pass validation inputs to network
- Compare the outputs (index with the highest probability) against the validation target
- Compute and append the validation accuracy



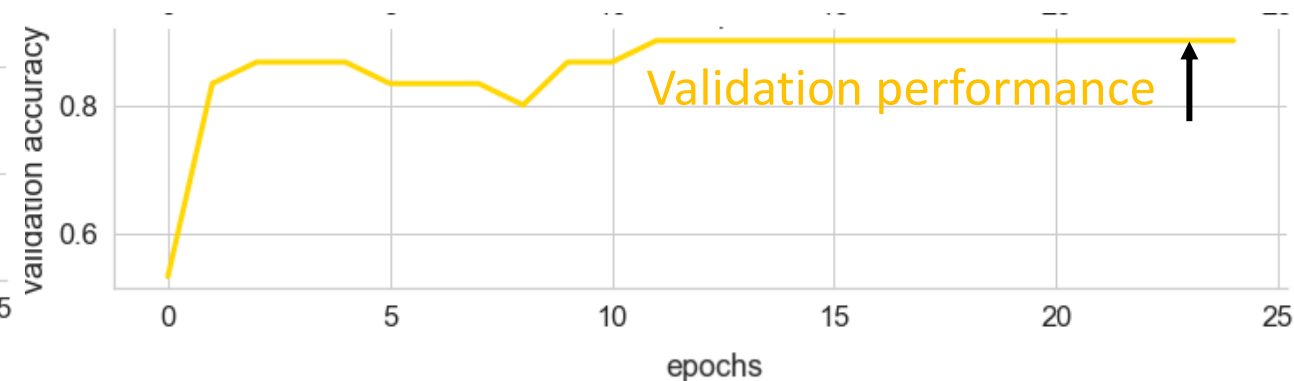


# Visualization and Evaluation

```
7 plt.subplot(2, 1, 2)
8 plt.plot(train_loss_list, linewidth = 3)
9 plt.ylabel("training loss")
10 plt.xlabel("epochs")
11 sns.despine()
```



```
1 plt.figure(figsize = (12, 6))
2
3 plt.subplot(2, 1, 1)
4 plt.plot(validation_accuracy_list, linewidth = 3, color = 'gold')
5 plt.ylabel("validation accuracy")
```



```
1 with torch.no_grad():
2
3     # Pass the testing feature data (30 samples) to the network to produce model predictions
4     y_pred_test = model(testing_inputs)
5
6     # Use the same technique as above to compute the testing classification accuracy
7     correct = (torch.argmax(y_pred_test, dim=1) == testing_targets).type(torch.FloatTensor)
8
9     print("Testing Accuracy: " + str(correct.mean().numpy()*100) + '%')
```

Testing Accuracy: 93.33333373069763%

93% classification accuracy

Testing performance ↑

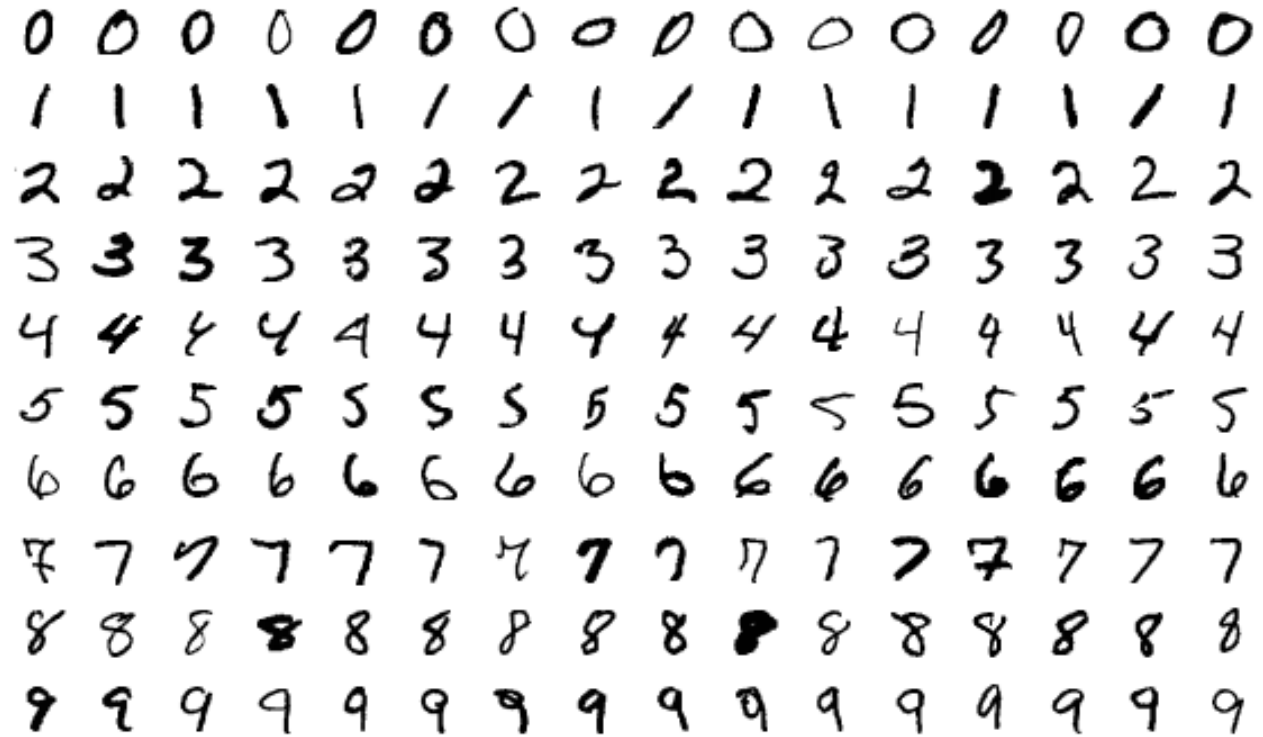


# LAB 3 ASSIGNMENT:

MNIST Classification using Fully Connected Network



# MNIST Dataset



Handwritten digits 0-9

Target labels are the correct values of the digit

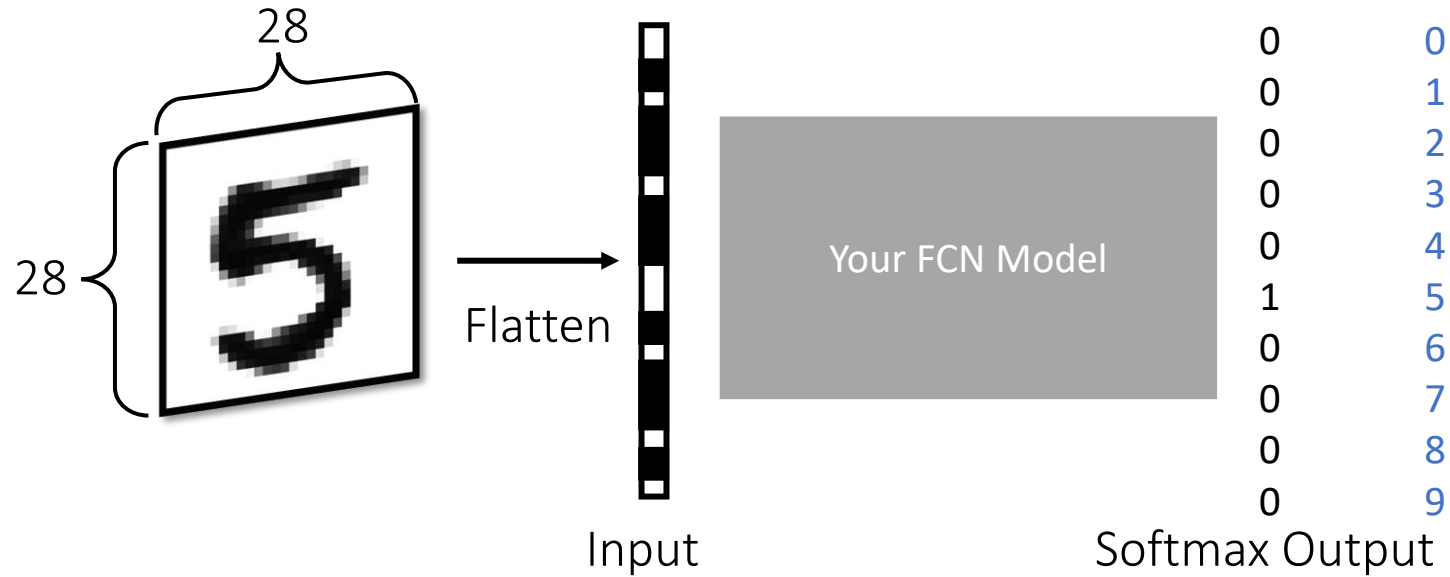
Data consists of grayscale images of fixed size (28x28) – flattens to 784

Canonical dataset for machine learning

1000 training samples, 100 testing samples



# MNIST Classification with FCN



In this exercise, you will classify handwritten digits (28 x 28) using your own **Fully Connected Network Architecture**.

Prior to training your neural net, 1) Flatten each digit into 1D array of size 784, 2) Normalize the dataset using standard scaler and 3) Split the dataset into train/validation/test.

Design your own neural net architecture with your choice of hidden layers, activation functions, optimization method etc.

Your goal is to **achieve a testing accuracy of >90%**, with no restrictions on epochs.

Demonstrate the performance of your model via plotting the **training loss, validation accuracy** and printing out the **testing accuracy**.

Plot the testing samples where your model failed to classify correctly and print your model's best guess for each of them



# Tips for Training Your Model

## First things to decide

- Number of layers
- Neurons in each layer
- Activation function  
(ReLU, Tanh, sigmoid)
- Training batch size  
(SGD, Mini-batch, Batch Gradient)
- Learning rate
- Optimizer  
(SGD, Adam, RMS Prop etc)
- Number of training epochs

## If your model is overfitting

(high training performance but low validation/testing performance)

- Add dropout layers
- Add regularization terms
- Stop training early
- Make network smaller (fewer layers or neurons)